

Machine Learning 03

Preprocessing - Categorical to number conversions

Claudio Sartori

1 Preprocessing: transform categorical data

In `scikit-learn` the classifiers require numeric data. The library makes available a set of preprocessing functions which help the transformation. This exercise proposes two types of transformations:

- `OneHotEncoder` for purely categorical columns: if the column has V distinct values it is substituted by V binary columns where in each row only the bit corresponding to the original value is true
- `OrdinalEncoder` for ordinal columns: the original V values are mapped into the $0..V-1$ range

The additional function `ColumnTransformer` allows to apply the different transformations to the appropriate columns with a single statement.

1.0.1 To do:

- import the appropriate names
- set the random state
- import the data set with the appropriate column names
- inspect the content and the data types
- read carefully the `.names` file of the data set, to understand which are the ordinal and categorical data
- data cleaning
 - the **ordinal transformer** generates a mapping from strings to numbers according to the lexicographic sorting of the strings; in this particular case, the strings indicate numeric subranges, and ranges with one digit constitute exceptions '5-9' happens to be after '20-25'
 - it is necessary to transform '5-9' into '05-09', and the same for other similar cases
 - a way to do this is to prepare dictionaries for the translation and use the `.map` function
- prepare the lists of the ordinal, categorical and numeric columns
- prepare the preprocessor
- split the cleaned data into the X and y part
- `fit_transform` the preprocessor and generate the transformed data set
- split the transformed data set into train and test
- use the same method used for the exercise of 19/11 to test several classifiers

<http://scikit->

learn.org/stable/auto_examples/model_selection/plot_grid_search_digits.html
 @author: scikit-learn.org and Claudio Sartori

```
[2]:
```

	Class	age	menopause	tumor-size	inv-nodes	node-caps	\
0	no-recurrence-events	30-39	premeno	30-34	0-2	no	
1	no-recurrence-events	40-49	premeno	20-24	0-2	no	
2	no-recurrence-events	40-49	premeno	20-24	0-2	no	
3	no-recurrence-events	60-69	ge40	15-19	0-2	no	
4	no-recurrence-events	40-49	premeno	0-4	0-2	no	

	deg-malig	breast	breast-quad	irradiat
0	3	left	left_low	no
1	2	right	right_up	no
2	2	left	left_low	no
3	2	right	left_up	no
4	2	right	right_low	no

Show the types of the columns

```
Class          object
age            object
menopause      object
tumor-size     object
inv-nodes      object
node-caps      object
deg-malig      int64
breast         object
breast-quad    object
irradiat       object
dtype: object
```

Clean the column tumor-size

```
[4]: {'30-34': '30-34',
      '20-24': '20-24',
      '15-19': '15-19',
      '0-4': '0-4',
      '25-29': '25-29',
      '50-54': '50-54',
      '10-14': '10-14',
      '40-44': '40-44',
      '35-39': '35-39',
      '5-9': '5-9',
      '45-49': '45-49'}
```

Clean the column inv-nodes

Inspect the data

```
[10]:
```

	Class	age	menopause	tumor-size	inv-nodes	node-caps	\
0	no-recurrence-events	30-39	premeno	30-34	00-02	no	
1	no-recurrence-events	40-49	premeno	20-24	00-02	no	
2	no-recurrence-events	40-49	premeno	20-24	00-02	no	
3	no-recurrence-events	60-69	ge40	15-19	00-02	no	
4	no-recurrence-events	40-49	premeno	00-04	00-02	no	

	deg-malig	breast	breast-quad	irradiat
0	3	left	left_low	no
1	2	right	right_up	no
2	2	left	left_low	no
3	2	right	left_up	no
4	2	right	right_low	no

Prepare the lists of numeric features, ordinal features, categorical features

The non-numeric features are:

```
['Class' 'age' 'menopause' 'tumor-size' 'inv-nodes' 'node-caps' 'breast'
 'breast-quad' 'irradiat']
```

The numeric features are:

```
['deg-malig']
```

The ordinal features are:

```
['age', 'tumor-size', 'inv-nodes']
```

The categorical features are:

```
['menopause', 'irradiat', 'breast', 'node-caps', 'breast-quad']
```

Prepare the transformer

Split X and y and check the shapes

The labels are:

```
['no-recurrence-events' 'recurrence-events']
```

```
[18]: (286, 9)
```

Fit the preprocessor with X and check the parameters printing the `.named_transformers_` attribute

```
[19]: ColumnTransformer(n_jobs=None, remainder='passthrough', sparse_threshold=0.3,
                        transformer_weights=None,
                        transformers=[('cat',
                                    OneHotEncoder(categorical_features=None,
                                                    categories=None, drop=None,
                                                    dtype=<class 'numpy.int32'>,
                                                    handle_unknown='ignore',
                                                    n_values=None, sparse=False),
                                    ['menopause', 'irradiat', 'breast',
                                    'node-caps', 'breast-quad'])),
```

```

('ord',
 OrdinalEncoder(categories='auto',
                 dtype=<class 'numpy.int32'>),
 ['age', 'tumor-size', 'inv-nodes']]),
verbose=False)

{'cat': OneHotEncoder(categorical_features=None, categories=None, drop=None,
                      dtype=<class 'numpy.int32'>, handle_unknown='ignore',
                      n_values=None, sparse=False), 'ord':
OrdinalEncoder(categories='auto', dtype=<class 'numpy.int32'>), 'remainder':
'passthrough'}

```

Fit-transform X and store the result in X_p, check the shape

[22]: (286, 20)

For ease of inspection transform X_p into a data frame df_p and inspect it

[24]:

	0	1	2	3	4	5 \
count	286.000000	286.000000	286.000000	286.000000	286.000000	286.000000
mean	0.451049	0.024476	0.524476	0.762238	0.237762	0.531469
std	0.498470	0.154791	0.500276	0.426459	0.426459	0.499883
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000
50%	0.000000	0.000000	1.000000	1.000000	0.000000	1.000000
75%	1.000000	0.000000	1.000000	1.000000	0.000000	1.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

	6	7	8	9	10	11 \
count	286.000000	286.000000	286.000000	286.000000	286.000000	286.000000
mean	0.468531	0.027972	0.776224	0.195804	0.003497	0.073427
std	0.499883	0.165182	0.417504	0.397514	0.059131	0.261293
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	1.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

	12	13	14	15	16	17 \
count	286.000000	286.000000	286.000000	286.000000	286.000000	286.000000
mean	0.384615	0.339161	0.083916	0.115385	2.664336	4.881119
std	0.487357	0.474254	0.277748	0.320046	1.011818	2.105930
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	2.000000	4.000000
50%	0.000000	0.000000	0.000000	0.000000	3.000000	5.000000
75%	1.000000	1.000000	0.000000	0.000000	3.000000	6.000000
max	1.000000	1.000000	1.000000	1.000000	5.000000	10.000000

	18	19
count	286.000000	286.000000
mean	0.517483	2.048951
std	1.110417	0.738217
min	0.000000	1.000000
25%	0.000000	2.000000
50%	0.000000	2.000000
75%	1.000000	3.000000
max	6.000000	3.000000

```
[25]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19
0	0	0	1	1	0	1	0	0	1	0	0	0	1	0	0	0	1	6	0	3
1	0	0	1	1	0	0	1	0	1	0	0	0	0	0	0	1	2	4	0	2
2	0	0	1	1	0	1	0	0	1	0	0	0	1	0	0	0	2	4	0	2
3	1	0	0	1	0	0	1	0	1	0	0	0	0	1	0	0	4	3	0	2
4	0	0	1	1	0	0	1	0	1	0	0	0	0	0	1	0	2	0	0	2

The columns in the transformed dataset are generated according to the order you see printing the preprocessor after fitting, therefore the last four columns correspond to 'age', 'tumor-size', 'inv-nodes', 'deg-malig'.

In order to inspect if the translation and check if the mapping is as expected, compare the sorted values of `df['tumor-size']` and `df_p[17]`, e.g. comparing the index sequences

The number of index discordances between 'tumor-size' and '17' is 0

Train/test split

Classification and test

```
=====
# Tuning hyper-parameters for recall_macro
```

```
-----
Trying model Decision Tree
Best parameters set found on train set:
```

```
{'max_depth': 14}
```

Grid scores on train set:

```
0.567 (+/-0.086) for {'max_depth': 1}
0.610 (+/-0.115) for {'max_depth': 2}
0.583 (+/-0.127) for {'max_depth': 3}
0.551 (+/-0.082) for {'max_depth': 4}
0.574 (+/-0.148) for {'max_depth': 5}
0.574 (+/-0.138) for {'max_depth': 6}
0.597 (+/-0.148) for {'max_depth': 7}
0.591 (+/-0.226) for {'max_depth': 8}
0.567 (+/-0.223) for {'max_depth': 9}
```

0.576 (+/-0.285) for {'max_depth': 10}
 0.577 (+/-0.168) for {'max_depth': 11}
 0.552 (+/-0.166) for {'max_depth': 12}
 0.564 (+/-0.163) for {'max_depth': 13}
 0.620 (+/-0.187) for {'max_depth': 14}
 0.565 (+/-0.142) for {'max_depth': 15}
 0.573 (+/-0.089) for {'max_depth': 16}
 0.608 (+/-0.121) for {'max_depth': 17}
 0.571 (+/-0.154) for {'max_depth': 18}
 0.576 (+/-0.177) for {'max_depth': 19}

Detailed classification report for the best parameter set:

The model is trained on the full train set.

The scores are computed on the full test set.

	precision	recall	f1-score	support
no-recurrence-events	0.72	0.84	0.77	49
recurrence-events	0.47	0.30	0.37	23
accuracy			0.67	72
macro avg	0.59	0.57	0.57	72
weighted avg	0.64	0.67	0.64	72

[[41 8]
 [16 7]]

 Trying model Gaussian Naive Bayes
 Best parameters set found on train set:

{'var_smoothing': 0.01}

Grid scores on train set:

0.500 (+/-0.000) for {'var_smoothing': 10}
 0.506 (+/-0.049) for {'var_smoothing': 1}
 0.593 (+/-0.115) for {'var_smoothing': 0.1}
 0.629 (+/-0.134) for {'var_smoothing': 0.01}
 0.627 (+/-0.125) for {'var_smoothing': 0.001}
 0.624 (+/-0.121) for {'var_smoothing': 0.0001}
 0.611 (+/-0.076) for {'var_smoothing': 1e-05}
 0.601 (+/-0.092) for {'var_smoothing': 1e-06}
 0.591 (+/-0.094) for {'var_smoothing': 1e-07}
 0.577 (+/-0.124) for {'var_smoothing': 1e-08}
 0.556 (+/-0.142) for {'var_smoothing': 1e-09}
 0.551 (+/-0.135) for {'var_smoothing': 1e-10}

Detailed classification report for the best parameter set:

The model is trained on the full train set.

The scores are computed on the full test set.

	precision	recall	f1-score	support
no-recurrence-events	0.73	0.88	0.80	49
recurrence-events	0.54	0.30	0.39	23
accuracy			0.69	72
macro avg	0.63	0.59	0.59	72
weighted avg	0.67	0.69	0.67	72

```
[[43  6]
 [16  7]]
```

Trying model Linear Perceptron

Best parameters set found on train set:

`{'early_stopping': True}`

Grid scores on train set:

0.564 (+/-0.111) for `{'early_stopping': True}`

Detailed classification report for the best parameter set:

The model is trained on the full train set.

The scores are computed on the full test set.

	precision	recall	f1-score	support
no-recurrence-events	1.00	0.14	0.25	49
recurrence-events	0.35	1.00	0.52	23
accuracy			0.42	72
macro avg	0.68	0.57	0.39	72
weighted avg	0.79	0.42	0.34	72

```
[[ 7 42]
 [ 0 23]]
```

Trying model Support Vector

Best parameters set found on train set:

```
{'C': 10, 'kernel': 'linear'}
```

Grid scores on train set:

```
0.500 (+/-0.000) for {'C': 1, 'gamma': 0.001, 'kernel': 'rbf'}
0.500 (+/-0.000) for {'C': 1, 'gamma': 0.0001, 'kernel': 'rbf'}
0.495 (+/-0.048) for {'C': 10, 'gamma': 0.001, 'kernel': 'rbf'}
0.500 (+/-0.000) for {'C': 10, 'gamma': 0.0001, 'kernel': 'rbf'}
0.549 (+/-0.064) for {'C': 100, 'gamma': 0.001, 'kernel': 'rbf'}
0.495 (+/-0.048) for {'C': 100, 'gamma': 0.0001, 'kernel': 'rbf'}
0.574 (+/-0.122) for {'C': 1000, 'gamma': 0.001, 'kernel': 'rbf'}
0.554 (+/-0.074) for {'C': 1000, 'gamma': 0.0001, 'kernel': 'rbf'}
0.582 (+/-0.091) for {'C': 1, 'kernel': 'linear'}
0.599 (+/-0.159) for {'C': 10, 'kernel': 'linear'}
0.599 (+/-0.159) for {'C': 100, 'kernel': 'linear'}
0.599 (+/-0.159) for {'C': 1000, 'kernel': 'linear'}
```

Detailed classification report for the best parameter set:

The model is trained on the full train set.

The scores are computed on the full test set.

	precision	recall	f1-score	support
no-recurrence-events	0.73	0.92	0.81	49
recurrence-events	0.60	0.26	0.36	23
accuracy			0.71	72
macro avg	0.66	0.59	0.59	72
weighted avg	0.69	0.71	0.67	72

```
[[45  4]
 [17  6]]
```

Trying model K Nearest Neighbor

Best parameters set found on train set:

```
{'metric': 'manhattan', 'n_neighbors': 7}
```

Grid scores on train set:

```
0.567 (+/-0.088) for {'metric': 'euclidean', 'n_neighbors': 1}
0.524 (+/-0.037) for {'metric': 'euclidean', 'n_neighbors': 2}
0.554 (+/-0.190) for {'metric': 'euclidean', 'n_neighbors': 3}
0.542 (+/-0.104) for {'metric': 'euclidean', 'n_neighbors': 4}
0.548 (+/-0.115) for {'metric': 'euclidean', 'n_neighbors': 5}
```


0.502 (+/-0.075) for {'metric': 'euclidean', 'n_neighbors': 6}
0.555 (+/-0.096) for {'metric': 'euclidean', 'n_neighbors': 7}
0.523 (+/-0.080) for {'metric': 'euclidean', 'n_neighbors': 8}
0.521 (+/-0.091) for {'metric': 'euclidean', 'n_neighbors': 9}
0.524 (+/-0.079) for {'metric': 'euclidean', 'n_neighbors': 10}
0.578 (+/-0.077) for {'metric': 'manhattan', 'n_neighbors': 1}
0.554 (+/-0.063) for {'metric': 'manhattan', 'n_neighbors': 2}
0.554 (+/-0.170) for {'metric': 'manhattan', 'n_neighbors': 3}
0.570 (+/-0.086) for {'metric': 'manhattan', 'n_neighbors': 4}
0.562 (+/-0.052) for {'metric': 'manhattan', 'n_neighbors': 5}
0.553 (+/-0.097) for {'metric': 'manhattan', 'n_neighbors': 6}
0.584 (+/-0.124) for {'metric': 'manhattan', 'n_neighbors': 7}
0.566 (+/-0.158) for {'metric': 'manhattan', 'n_neighbors': 8}
0.560 (+/-0.161) for {'metric': 'manhattan', 'n_neighbors': 9}
0.561 (+/-0.095) for {'metric': 'manhattan', 'n_neighbors': 10}
0.490 (+/-0.152) for {'metric': 'chebyshev', 'n_neighbors': 1}
0.521 (+/-0.128) for {'metric': 'chebyshev', 'n_neighbors': 2}
0.575 (+/-0.146) for {'metric': 'chebyshev', 'n_neighbors': 3}
0.539 (+/-0.090) for {'metric': 'chebyshev', 'n_neighbors': 4}
0.576 (+/-0.095) for {'metric': 'chebyshev', 'n_neighbors': 5}
0.518 (+/-0.087) for {'metric': 'chebyshev', 'n_neighbors': 6}
0.531 (+/-0.100) for {'metric': 'chebyshev', 'n_neighbors': 7}
0.518 (+/-0.068) for {'metric': 'chebyshev', 'n_neighbors': 8}
0.520 (+/-0.086) for {'metric': 'chebyshev', 'n_neighbors': 9}
0.539 (+/-0.070) for {'metric': 'chebyshev', 'n_neighbors': 10}

Detailed classification report for the best parameter set:

The model is trained on the full train set.

The scores are computed on the full test set.

	precision	recall	f1-score	support
no-recurrence-events	0.69	0.92	0.79	49
recurrence-events	0.43	0.13	0.20	23
accuracy			0.67	72
macro avg	0.56	0.52	0.49	72
weighted avg	0.61	0.67	0.60	72

```
[[45  4]
 [20  3]]
```

Trying model Random Forest

Best parameters set found on train set:

```
{'max_depth': 8}
```

Grid scores on train set:

0.533 (+/-0.082) for {'max_depth': 1}
0.604 (+/-0.076) for {'max_depth': 2}
0.587 (+/-0.130) for {'max_depth': 3}
0.595 (+/-0.111) for {'max_depth': 4}
0.599 (+/-0.194) for {'max_depth': 5}
0.590 (+/-0.115) for {'max_depth': 6}
0.580 (+/-0.117) for {'max_depth': 7}
0.614 (+/-0.110) for {'max_depth': 8}
0.560 (+/-0.087) for {'max_depth': 9}
0.600 (+/-0.112) for {'max_depth': 10}

Detailed classification report for the best parameter set:

The model is trained on the full train set.

The scores are computed on the full test set.

	precision	recall	f1-score	support
no-recurrence-events	0.73	0.96	0.83	49
recurrence-events	0.75	0.26	0.39	23
accuracy			0.74	72
macro avg	0.74	0.61	0.61	72
weighted avg	0.74	0.74	0.69	72

[[47 2]
[17 6]]

Trying model Adaboost

Best parameters set found on train set:

{'learning_rate': 0.01}

Grid scores on train set:

0.586 (+/-0.146) for {'learning_rate': 1.0}
0.620 (+/-0.102) for {'learning_rate': 0.1}
0.643 (+/-0.161) for {'learning_rate': 0.01}
0.567 (+/-0.086) for {'learning_rate': 0.001}
0.567 (+/-0.086) for {'learning_rate': 0.0001}

Detailed classification report for the best parameter set:

The model is trained on the full train set.

The scores are computed on the full test set.

	precision	recall	f1-score	support
no-recurrence-events	0.72	0.98	0.83	49
recurrence-events	0.80	0.17	0.29	23
accuracy			0.72	72
macro avg	0.76	0.58	0.56	72
weighted avg	0.74	0.72	0.65	72

```
[[48  1]
 [19  4]]
```

Summary of results for recall_macro

Estimator

Decision Tree	- score: 0.62%
Gaussian Naive Bayes	- score: 0.63%
Linear Perceptron	- score: 0.56%
Support Vector	- score: 0.6%
K Nearest Neighbor	- score: 0.58%
Random Forest	- score: 0.61%
Adaboost	- score: 0.64%